

Spotify Recommendation Platform

(Anand Rastogi – 2022OG04054)

Problem: Spotify Music Recommendation Platform

With multiple audio features such as danceability, energy, and valence, this dataset can be used to build a music recommendation system. By analyzing the preferences of users for certain genres and characteristics of tracks, the system can suggest similar tracks or even recommend new genres that users might enjoy.

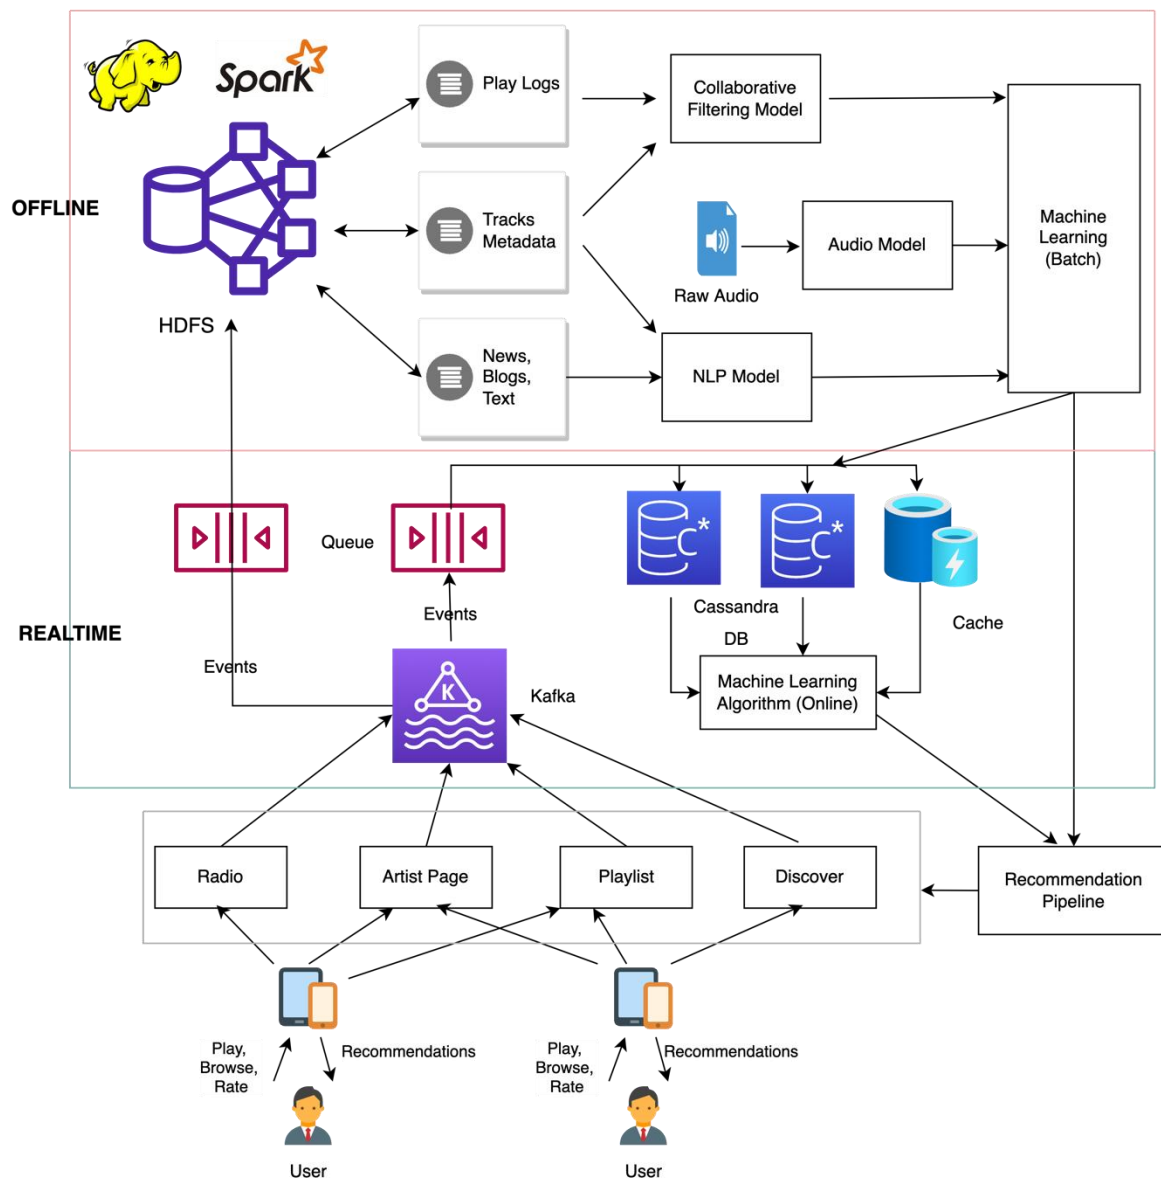
End-User experience: As a user searches or plays a particular song in Spotify, your Music Recommendation Platform will be able to return other music recommendations for the user. And if the user selects or ignores the recommendation, that data is fed back to improve the recommendation algorithm.

1) **Design:** Design the backend services, APIs, data models and the databases you will use for this. The design should support both real-time stream processing and batch analytics use cases. Explain the rationale for your design especially tradeoffs w.r.t CAP theorem, describe the tech choices for the databases and other components. You must use a NoSQL database as part of your design (MongoDB, Neo4j, DynamoDB, HBase, anything appropriate for your solution).

Solution notes: Submission for this part MUST include an Architecture diagram with bullet points or a brief description of the design choices, rationale, tech choices.

GitHub Link -

https://github.com/AndyBITS/BDS_Assignment_SpotifyRecommendations.git



- 1) Spotify Music Recommendation Platforms will comprise processing of real time events in Realtime as well as Offline mode.
- 2) User plays, browses and rates the songs from Spotify access points such as Apps, and WebUI etc. Every song browsed, played or skipped would generate an event which will be sent through Kafka streaming layer.
- 3) Realtime processing of events will enhance the user's experience by providing the recommendations based on recent activities such as Continue Playing, Recently Listened Songs, Artists Playlist and Genre.

- 4) Offline processing of events will enhance the user's experience by providing the recommendations from ML Models implemented on aggregated data over a long period.
 - a. Collaborative Filtering Model: Recommendations based on similarity and behavior identified among users.
 - b. NLP Models: Recommendations on news articles and blogs published for artists, music tracks and albums.
 - c. Content Filtering Model: Recommendations based on insights gained by analysis of songs metadata.
 - d. Audio Filtering: Recommendations based on raw audio analysis to get the insights on danceability, energy and valence.

Infrastructure Design/Proposal

- 1) Kafka is used for streaming the events.
- 2) Cassandra (NoSQL) DB is proposed for real-time processing mode.
 - a. Aligns with Availability and Partition Tolerance principles in CAP Theorem.
 - b. Scalability: It ensures highly availability by replicating data within and across the nodes.
 - c. Read Optimized: Cassandra is based on Query First Approach, and we can design the tables based on the queries we are going to design. The data will be loaded in partitions spread across the nodes, enabling the fast retrieval. E.g.
 - i. The data with different genres, artists will be loaded in different nodes and help with fast retrievals with simple queries (not requiring join operations).
 - ii. The write/update on specific genre will not lock the partitions with another genre.
 - d. Eventual Consistency: The data consistency requirement is not immediate in music recommendation platform and we should be good with near time latency for data to be updated across replicas in all nodes.
 - e. Cache has also been enabled to handle the read/write latency from ML algorithms.
 - f. Cassandra is scalable and we can add more nodes whenever we need more data to be stored.
- 3) Hadoop Cluster is proposed for offline processing mode to store the logs, events and existing metadata for tracks. Spark can be used for processing.

4) Machine Learning Algorithms can be implemented using implemented through Spark.

2) **OLTP Queries:** Execute at least 3 key read/write queries against the NoSQL database as part of the solution with the optimizations and the consistency levels set in your queries.

Solution notes: Submission for this part MUST include a pdf of the queries and a video of their successful execution (you can use loom for the video). No more than 2 pages for the queries and no more than 1 minute for the loom video.

Since we are designing the music recommendation platform and Cassandra follows a query first approach model, our select queries will be mainly used for retrieving and recommending the music tracks on features such as track_genre, artists, and popularity etc. Hence, we should have the tables created with these attributes as the partition keys for optimized results.

Since Cassandra was installed with one node on personal laptop, I could not capture the instances with consistency levels other than One. I did try to change level to Any/Quorum once but I started getting error notifications.

Recording link -

<https://www.loom.com/share/e3a116ac1fc04341ad90698f3180cb7d?sid=c6e9457d-1359-4a8d-907b-ee94aca6e6d3>

Below are the sample READ/WRITE queries we will execute as part of this assignment:

a. **Retrieve the 10 most popular songs for a specific artist.**

```
cqlsh:music_keyspace> select track_name, popularity, artists from
music_by_artists_popularity where artists = 'Acid Ghost' order by
popularity desc limit 10;
```

track_name	popularity	artists
Sandy Kim	57	Acid Ghost
I Don't Need You	53	Acid Ghost
New York	52	Acid Ghost
Overthinking	50	Acid Ghost
All Alone	49	Acid Ghost
Not the Same Person	49	Acid Ghost
Nature (Poem)	48	Acid Ghost
Hide My Face	45	Acid Ghost

(8 rows)
 cqlsh:music_keyspace>

b. Retrieve any 10 songs, album and artist for a specific track_genre.

cqlsh:music_keyspace> select track_name, album_name, artists from music_by_track_genre where track_genre ='rock' limit 10;

track_name	album_name	artists
Mr. Brightside	The Killers	Hot Fuss
Merry Christmas	Bryan Adams	Christmas Dinner 2022
Little Saint Nick - 1991 Remix	The Beach Boys	Best Christmas Playlist Ever 2022
I Put A Spell On You	Creedence Clearwater Revival	Pumpkin Patch Hits
Summer Of '69	Bryan Adams	Best Of Karneval
All My Favorite Songs	Weezer;AJR	Here to Forever - All New Indie
There'd Better Be A Mirrorball	Arctic Monkeys	I Ain't Quite Where I Think I Am
Christmas Without You	OneRepublic	CHRISTMAS TOP HITS 2022
Centuries	Fall Out Boy	American Beauty/American Psycho
Blood // Water	grandson	All Out Alternative

c. Insert a new track for a new artist and retrieve the row to confirm.

cqlsh:music_keyspace> consistency one
 Consistency level set to ONE.

cqlsh:music_keyspace> insert into music_by_artists_popularity (artists, popularity, track_id, album_name, track_name, duration_ms, explicit, danceability, energy, key, loudness, mode, speechiness, acousticness, instrumentalness, liveness, valence, tempo, time_signature, track_genre) VALUES ('Ghulam Ali',100,'1ne2w0tXF2jM6Cm8WBkaER','Hungama','Hungama Hein Kyun Barpa',55555, FALSE,0.421,0.637,3,-1000,1,0.01,0.1,0,1,1,100,4,'Retro Ghazal');

cqlsh:music_keyspace> select * from music_by_artists_popularity where artists = 'Ghulam Ali';

artists	popularity	track_id	album_name	danceability	duration_ms	energy	explicit	instrumentalness	key	loudness	mode	speechiness	acousticness	liveness	valence	tempo	time_signature	track_genre
---------	------------	----------	------------	--------------	-------------	--------	----------	------------------	-----	----------	------	-------------	--------------	----------	---------	-------	----------------	-------------

	0.147	147.591		4	black-metal	Riders Of
The Apocalypse		0.236				
1349		20	0Js35d0JcXFyXnVQgZrYx1		2.3e-06	
		Hellfire	0.169	465586	0.981	
False		0.788	0	0.303	-5.07	1
	0.133	117.6		4	Ghazals	Celestial
Deconstruction		0.0393				
1349		20	60edNIKZXc0i95A3BcNIms		4.3e-06	The
Infernal Pathway		0.0998	330066	0.911	False	
	0.714	9	0.18	-11.959	0	0.112
195.039		4	black-metal		Abyssos	
Antithesis		0.0708				

(5 rows)

3) **OLAP Analytics Queries:** Within the broader context of this problem (Spotify) you need to perform analytics queries from the dataset to driver better Sales / Growth. For this, setup a big data batch processing environment (Hadoop or a Spark), share the config. Then perform at least 3 queries in the form of Map-Reduce jobs or Pig / HiveQL / SparkSQL queries and share the responses.

Solution notes: Submission for this part MUST include a pdf of the environment configuration, the queries and a video of their successful execution (you can use loom for the video). No more than 2 pages for the config and queries and no more than 1 minute for the loom video.

Refer the files available at GitHub link - .

https://github.com/AndyBITS/BDS_Assignment_SpotifyRecommendations/blob/main/BDS%20Assignment_Spotify%20recommendations_Config_Anand%20Rastogi_2022OG04054.html

https://github.com/AndyBITS/BDS_Assignment_SpotifyRecommendations/blob/main/BDS%20Assignment_Spotify%20recommendations_OLAP%20Analytics_Anand%20Rastogi_2022OG04054.html



BDS
Assignment_Spotify



BDS
Assignment_Spotify

4) Appendix

DDL/ DML SQL

```
cqlsh> CREATE KEYSPACE music_keyspace WITH replication =
{'class':'SimpleStrategy','replication_factor':3} AND
durable_writes = 'true';
```

Warnings :

Your replication factor 3 for keyspace music_keyspace is higher than the number of nodes 1.

```
cqlsh> USE music_keyspace;
cqlsh:music_keyspace>
```

```
cqlsh:music_keyspace> CREATE TABLE music_by_track_id (track_id
text PRIMARY KEY, artists text, album_name text, track_name
text, popularity int, duration_ms int, explicit boolean,
danceability float, energy float, key int, loudness float, mode
int, speechiness float, acousticness float, instrumentalness
float, liveness float, valence float, tempo float, time_signature
int, track_genre text);
cqlsh:music_keyspace> describe table music_by_track_id
```

```
CREATE TABLE music_keyspace.music_by_track_id (
    track_id text PRIMARY KEY,
    acousticness float,
    album_name text,
    artists text,
    danceability float,
    duration_ms int,
    energy float,
    explicit boolean,
    instrumentalness float,
    key int,
    liveness float,
    loudness float,
    mode int,
    popularity int,
    speechiness float,
    tempo float,
    time_signature int,
    track_genre text,
    track_name text,
    valence float
) WITH additional_write_policy = '99p'
    AND bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND cdc = false
    AND comment = ''
```



```

        AND compaction = {'class':
'org.apache.cassandra.db.compaction.SizeTieredCompactionStrate
gy', 'max_threshold': '32', 'min_threshold': '4'}
        AND compression = {'chunk_length_in_kb': '16', 'class':
'org.apache.cassandra.io.compress.LZ4Compressor'}
        AND memtable = 'default'
        AND crc_check_chance = 1.0
        AND default_time_to_live = 0
        AND extensions = {}
        AND gc_grace_seconds = 864000
        AND max_index_interval = 2048
        AND memtable_flush_period_in_ms = 0
        AND min_index_interval = 128
        AND read_repair = 'BLOCKING'
        AND speculative_retry = '99p';

```

```
cqlsh:music_keyspace> select count(*) from music_by_track_id;
```

```

count
-----
89621

```

```
(1 rows)
```

Warnings :

Aggregation query used without partition key

```

cqlsh:music_keyspace> CREATE TABLE music_by_artists_popularity
(artists text, popularity int, track_id text , album_name text,
track_name text, duration_ms int, explicit boolean,
danceability float, energy float, key int, loudness float, mode
int, speechiness float, acousticness float, instrumentalness
float, liveness float, valence float, tempo float, time_signature
int, track_genre text, PRIMARY KEY (artists,
popularity, track_id));
cqlsh:music_keyspace> describe table
music_by_artists_popularity

```

```

CREATE TABLE music_keyspace.music_by_artists_popularity (
    artists text,
    popularity int,
    track_id text,
    acousticness float,
    album_name text,
    danceability float,
    duration_ms int,
    energy float,
    explicit boolean,
    instrumentalness float,

```

```

    key int,
    liveness float,
    loudness float,
    mode int,
    speechiness float,
    tempo float,
    time_signature int,
    track_genre text,
    track_name text,
    valence float,
    PRIMARY KEY (artists, popularity, track_id)
) WITH CLUSTERING ORDER BY (popularity ASC, track_id ASC)
    AND additional_write_policy = '99p'
    AND bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND cdc = false
    AND comment = ''
    AND compaction = {'class':
'org.apache.cassandra.db.compaction.SizeTieredCompactionStrate
gy', 'max_threshold': '32', 'min_threshold': '4'}
    AND compression = {'chunk_length_in_kb': '16', 'class':
'org.apache.cassandra.io.compress.LZ4Compressor'}
    AND memtable = 'default'
    AND crc_check_chance = 1.0
    AND default_time_to_live = 0
    AND extensions = {}
    AND gc_grace_seconds = 864000
    AND max_index_interval = 2048
    AND memtable_flush_period_in_ms = 0
    AND min_index_interval = 128
    AND read_repair = 'BLOCKING'
    AND speculative_retry = '99p';

```

```

cqlsh:music_keyspace> select count(*) from
music_by_artists_popularity;

```

```

count
-----
90340

```

```

(1 rows)

```

```

Warnings :
Aggregation query used without partition key

```

```

cqlsh:music_keyspace>

```

```
cqlsh:music_keyspace> CREATE TABLE music_by_track_genre
(track_genre text, track_id text, artists text, album_name text,
track_name text, popularity int, duration_ms int, explicit
boolean, danceability float, energy float, key int, loudness
float, mode int, speechiness float, acousticness
float, instrumentalness float, liveness float, valence
float, tempo float, time_signature int, PRIMARY KEY
(track_genre, track_id));
cqlsh:music_keyspace> describe music_by_track_genre
```

```
CREATE TABLE music_keyspace.music_by_track_genre (
    track_genre text,
    track_id text,
    acousticness float,
    album_name text,
    artists text,
    danceability float,
    duration_ms int,
    energy float,
    explicit boolean,
    instrumentalness float,
    key int,
    liveness float,
    loudness float,
    mode int,
    popularity int,
    speechiness float,
    tempo float,
    time_signature int,
    track_name text,
    valence float,
    PRIMARY KEY (track_genre, track_id)
) WITH CLUSTERING ORDER BY (track_id ASC)
    AND additional_write_policy = '99p'
    AND bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND cdc = false
    AND comment = ''
    AND compaction = {'class':
'org.apache.cassandra.db.compaction.SizeTieredCompactionStrate
gy', 'max_threshold': '32', 'min_threshold': '4'}
    AND compression = {'chunk_length_in_kb': '16', 'class':
'org.apache.cassandra.io.compress.LZ4Compressor'}
    AND memtable = 'default'
    AND crc_check_chance = 1.0
    AND default_time_to_live = 0
    AND extensions = {}
    AND gc_grace_seconds = 864000
    AND max_index_interval = 2048
    AND memtable_flush_period_in_ms = 0
    AND min_index_interval = 128
```

```
    AND read_repair = 'BLOCKING'
    AND speculative_retry = '99p';
cqlsh:music_keyspace>
```

```
cqlsh:music_keyspace> select count(*) from
music_by_track_genre;
```

```
count
```

```
-----
```

```
113421
```

```
(1 rows)
```

```
Warnings :
```

```
Aggregation query used without partition key
```

```
cqlsh:music_keyspace>
```